

Briefing for Thursday’s Call with NG

The Ten Highlighted Exhibits: Methods, Likely Questions, Honest Assessment

July 22, 2026

Numbering mapped to the current v1.1. For each item: what it is, exactly how it’s built (data → computation → script), what NG is likely asking, and an honest assessment.

The shape of his list. The ten highlights are really three method blocks:

1. **The calibration** — Table 4.
2. **The LLM-similarity module** — Tables 5, 6, 7, Figure 6, and the “From contexts to categories” paragraph (his “I do not understand”).
3. **The heterogeneity accounting** — eq. (7), Table 13 (Market row), Figure 13, Table 14.

His two margin comments (“beliefs about P2?” on eq. 7; “why mean? are beliefs changing?” on Figure 13) are the *same* question asked twice: **where does the belief component of the representation live in each accounting exercise?** The answer — some exercises deliberately use only the category (attention) component, others add beliefs — is the single most useful thing to have crisp on the call. A one-slide map would defuse both notes:

Exercise	Category shares	Beliefs	Nature
Fig. 12 (fig:instability_all)	yes	yes (+ fitted P1 action for P2)	regression-based
Eq. 7 / Table 13 (BCS)	yes	no	variance decomposition
Fig. 13 (mean & SD reweighting)	yes	no	regression-free counterfactual
Table 14 (Oaxaca)	yes	yes (terciles, UG/TG cells)	regression-free, preregistered; bootstrap SEs

1 Table 4 — Calibrated Representations, by Category (Control Sample)

Where it lives. Script 08_calibration.py in replication_package/code/, writing output/tables/calibration_p1.tex; full log in calibration_stats.txt alongside. Input: data/player1_all_categorized.xlsx, control condition only, classified P1 responses only (three substantive categories).

How it works — the recursive chain. All moments are category-level means; each participant plays one game, so moments for the same category come from different people:

1. **DG pins σ/μ :** $\sigma/\mu = \max(0, 1/2 - \bar{t}_{\text{DG}})$ (the interior DG first-order condition). Moral: $\bar{t} = .462 \rightarrow .038$. Self-interest: $\bar{t} = .103 \rightarrow .397$.

2. **TG pins ρ/μ :** the TG FOC at the measured mean believed share s (hypothetical belief at the 1/3 reference), *including the equal-payoff anchor* $t^e(s) = 1/[1 + (1+r)(1-2s)]$: $\rho/\mu = [\bar{r}_{\text{TG}} - t^e(s) + (\sigma/\mu)(1+r)(1-s)]/r$, with $r = 2$.
3. **UG pins (a,b) :** two equations — the reference belief point $\bar{p} = a + b/3$ and the UG FOC at the category's mean offer, linear in b . For **Moral** the solution has $b \leq 0$: no valid schedule, printed “–”. This is a *feature*: $\sigma/\mu = .038$ means behavior is norm-dominated and insensitive to beliefs, so beliefs are unidentified exactly where the model says they should be.
4. **MBC** has a negligible DG cell (~ 3 obs, 0.8%), so no DG moment: the Self-interest schedule (a,b) is imposed (the $\dagger$ in the table; justified by reference-action acceptance beliefs being flat across categories, 48.2–48.5%), and $(\sigma/\mu, \rho/\mu)$ solve the UG FOC + TG locus jointly (Brent root finder on $\sigma/\mu \in [0, .499]$).
5. **Bootstrap:** participants resampled with replacement within game, $B = 1,000$ (all successful), re-running the whole inversion; SEs = std of draws.

The overidentifying test (last two columns): the calibration never uses the belief each proposer attaches to her *chosen* offer. Implied acceptance at the category's mean chosen offer (from the calibrated schedule) vs the measured chosen-action belief: .565 vs .795 (Moral), .525 vs .701 (SI), .598 vs .711 (MBC) — the +11 to +23pp chosen-action optimism the forecast-error section picks up.

What NG likely wants. The mechanics above, plus the identification logic: the §3.2 paragraph added on 2026-07-20 already states that category-level pooling is *forced* by the between-subject design (each moment comes from different people), and that this is the model's own claim — attention attaches to the situation, not the person.

Honest assessment.

- *Sound, provided it is presented as a calibration, not an estimation.* It inverts interior FOCs at category *means* while the actual distributions are bimodal (mass at 0 and at the equal split); the “representative agent per category” is a portrait. The defense is exactly the paper's: three non-imposed regularities ($\rho \approx \sigma$ for SI; Moral norm-dominated with beliefs unidentified exactly where the model says; $a + b = .81 \leq 1$ a genuine probability schedule) plus the overid test.
- *Two-point route:* if NG asks “why not use both belief points per proposer to get (a,b) ?” — we tried; it's in the script as a diagnostic: it yields $a < 0$ and $a + b \approx 1.7$ –2.6, invalid, precisely because of chosen-action optimism. Decision: FOC route is the headline, two-point kept as diagnostic. This is a *strength* to volunteer, not hide.
- *Small gap — fixed (2026-07-20):* the table note now says that Moral's “implied” column uses the Self-interest schedule, its own being unidentified (the code falls back to it when Moral's schedule is invalid). Fixed in the generator and regenerated; all numbers byte-identical (seeded bootstrap).
- MBC's $\sigma/\mu = 0.000$ is a boundary solution (bootstrap SE .044 spans it) — have that ready if he asks how attention to own earnings can be *exactly* zero.

2 Table 5 — Perceived Similarity of Stories to Games (Round 1)

Where it lives. Until 2026-07-20 hand-typed in `main.tex` from the round-1 workbook `memory_games_llm_recording.xlsx` in `llm_similarity/`; nine transcripts alongside it. **Gap closed:** 21_round1_

`similarity_table.py` now recomputes every cell from the raw rating sheets (72 similarity ratings, 36 retrieval splits), asserts the means against the workbook’s Summary sheet, and emits `output/tables/llm_similarity.tex`, verified byte-identical to the hand-typed block; `main.tex` now inputs it, and the workbook is copied into `replication_package/data/`.

How it works. Three frontier models (GPT-5.5 Thinking, Claude Opus 4.7, Gemini 3 Pro) $\times$ three conversations each with independently permuted neutral labels (games A–D, Story 1/Story 2 — masking both position and priors about the experimental literature). Each conversation: (i) story–game similarity 0–100 for all 8 pairs ($\rightarrow$ 72 ratings over 9 conversations), (ii) a five-dimension structural decomposition of the games, (iii) a 100-point “which story does this game call to mind” retrieval split. Table cells are means over the nine conversations.

What it’s for. Three findings, all *ordinal*: (a) structural distance of stories to games is sharply ordered DG-KW 94.6 > DG-LT 51.2 > UG 29.9 > TG 14.4 (every model isolates the TG via the expandable pie); (b) **Bonus and Aid are structurally indistinguishable** (≤ 1.1 apart in every game) — so story effects cannot run through perceived game structure, which is what the F_c lever requires; (c) the retrieval split diverges from structural similarity (Bonus share rises 41 $\rightarrow$ 72% with strategic richness) — the structure-vs-context wedge the retrieval model runs on.

Honest assessment.

- The obvious objection — LLM raters are not our participants — is conceded in the paper; the pre-registered human survey replicating both rounds is the validation path. What makes the exhibit usable *now* is (i) three independent model families agreeing in every conversation, (ii) only ordinal/contrast claims are consumed downstream (the Table-7 gradient and the Bonus $\approx$ Aid equivalence), never the cardinal levels.
- Minor disclosure point to have ready: the workbook log records that Gemini’s *structural-decomposition* prompt in triplets 1–2 was answered by Gemini 3.5 Flash (capacity fallback). The similarity/retrieval numbers in Table 5 are not affected, but if we ever enumerate raters task-by-task, this belongs in the note.

3 The “From contexts to categories” paragraph (NG: “I do not understand”)

What it says, in plain language (this is the walk-through for the call):

1. The retrieval equation needs a similarity input $S(\text{context}, \text{category})$. A category, as §2.2 defines it, is a *class of concrete everyday situations*, not a word. You cannot ask a rater “how similar is this text to Morality?” without measuring the label instead of the class.
2. So we made each category concrete: **8 short vignettes per P1 category** (24 total) — everyday situations instantiating the category definitions the classifier uses — written by one model (Claude Fable 5) in a single frozen conversation, with banned category-naming words and the same eight life domains per category (so category $\neq$ setting). Receiver types (accept/reject, return/keep) get 2 vignettes each (32 vignettes total).
3. A *different* model (Claude Opus 4.8, reasoning on), in fresh conversations, **blind to everything** (vignettes = V1–V24, contexts = TEXT A–J, three permuted orders), rates each of the 10 context texts — the two stories plus the verbatim Control and Market instructions of the four games — against each

vignette (0–100), then distributes 100 “which would come to mind” retrieval points, and for UG/TG splits 100 points across the receiver vignettes.

4. **The rater never sees a category.** Category similarity is computed only afterwards, by mapping each vignette back to the category it instantiates and averaging its eight vignettes. That is the sentence NG most likely tripped on — it’s the design’s blinding device, not an afterthought.

Why this design (the answer to “is this the right approach?”): it measures the theory’s own object (similarity to situation classes); blinding kills demand effects and hypothesis-tailoring; ex-post mapping means the rater physically cannot know what is being tested. Circularity worry (vignettes come from the same category definitions as the classifier): intentional — both instruments measure the same theoretical construct; the two measurements are otherwise fully independent (different objects: instruction texts vs participant free text; blind raters), so nothing mechanically forces Figure 6’s correlation.

Suggestion: on the call, do one concrete example (e.g., the TG Market text against one cooperation vignette and one moral vignette). If NG still finds the paragraph opaque, offer to add a small design-schema figure or split the paragraph in two (generation / rating) — easy post-call edit.

4 Table 6 — Similarity of Contexts to Representation Categories (Round 2)

Where it lives. `llm_similarity/round2/04_analysis.py` → `replication_package/output/tables/similarity_categories.tex`. Raw ratings in `round2/recording_round2.csv` + `recording_splits_round2.csv`; transcripts and freeze record alongside. Filtered to the **Opus 4.8 headline rater only** (per protocol; Fable 5 enters robustness checks only).

How each cell is computed.

- *Similarity block:* mean of the 0–100 ratings over that category’s 8 vignettes × 3 permuted conversations (24 ratings per cell).
- *Retrieval-split block:* the category’s share of the 100 retrieval points, pooled over the 3 conversations.
- The receiver-type splits quoted in the text (accepting 42.7→52.0 under UG Market; TG 53.0→54.0) come from the same module and are now **Table 40** (tab:receiver_splits, Appendix H — added 2026-07-22, generated from the raw recording by `round2/07_receiver_splits_table.py`; Appendix H follows every existing float, so no existing exhibit number changed).

The three findings consumed in the text: Market–Control moral-similarity drops (–34.2/–29.5/–25.0 in DG-KW/DG-LT/UG) with SI and MBC both rising — matching the measured representation shifts; Aid vs Bonus nearly identical on Moral but split on SI (+2.5 toward Aid) and MBC (+2.1 toward Bonus); the control TG is the cooperation-similar outlier (47.3 vs 9–13) and the least moral-similar (12.3) — the similarity counterpart of Figure 2.

Honest assessment.

- *Single-model working measurement* — the paper says so explicitly, and the robustness appendix (H) is strong for what it is: test–retest $r = .994$ across the full re-run; second rater from a different tier (Fable 5) $r = .982$ with the refusal saga disclosed verbatim; permutation ranges reported (median 12.6 points across conversations — the *three-conversation means* are what’s stable, not single conversations).

- The genuine soft spot: **generator and both raters are one model family** (Anthropic). Disclosed in the appendix, and the multi-model grid (needs GPT/Gemini keys) + human survey Module B are the stated path. If NG asks “would GPT reproduce this?” the honest answer is: that is exactly the next step, wired and waiting on keys; no paper claim currently rests on the module alone.
- Retrieval-split columns are winner-take-most (Moral 89–94% in fixed-pie control contexts) — expected from a forced 100-point split over 24 vignettes; the quoted shifts and Figure 6 use the graded similarity block, not the splits.

5 Figure 6 — Representation Shifts vs Similarity Shifts (H5)

Where it lives. Same script (round2/04_analysis.py) → similarity_h5.pdf; stats in round2/out/h5_stats.txt.

How it’s built. 18 points = 3 categories × 3 games (DG-KW, UG, TG) × 2 comparisons. y = change in the *classified* category share from the experiment (Market–Control, Aid–Bonus; share in fractions). x = change in Table 6’s similarity column (Market–Control: the two instruction texts, game-specific; Aid–Bonus: the two story texts — so the 9 Aid–Bonus points share **three** x -values, identical across games). OLS with intercept, unweighted: slope .0122 (se .0025), $R^2 = .59$, Spearman $\rho = .57$. DG-LT excluded (its classified-share treatment effects aren’t part of the §4 figures); M1 leave-one-out held out per SN’s decision (slope .011, $R^2 = .50$ — restorable).

Honest assessment — know these three limitations before he finds them.

1. **The 18 points are not independent:** within each game×comparison triplet the y ’s sum to exactly zero (shares sum to one), and the cells share samples. The paper treats the fit as descriptive; the reported se is indicative. Say that proactively.
2. **The Aid–Bonus x ’s are tiny and game-invariant** (+2.6/+2.5/−2.1): those 9 points mostly sit near the origin; the slope is driven by the Market–Control fixed-pie cells. That is fine — the market frame is where similarity moves — but the figure should not be oversold as 18 independent confirmations.
3. **The TG Market–Control cells are off-pattern by design of the world, not of the test:** similarity barely moves ($\leq +1.5$) while representations and behavior move a lot. The paper’s reading (the TG instruction text already reads as a joint venture, so the market overlay adds no *resemblance*; its effects must run through availability F_c or beliefs) is candid and flagged as the sharpest target for the human survey. This is the right answer for NG: the exceptions localize the levers, and they are testable.

Net: correct approach *as a coherence check across two independent measurement systems*; it would be the wrong approach if presented as an estimated structural elasticity — the text already frames it correctly.

6 Table 7 — Story Effects Track Structural Similarity

Where it lives. 05_paper_v2_new_outputs.py (block E5, in replication_package/code/) → similarity_gradient.tex. Similarity and retrieval columns are hard-coded from the audited round-1 workbook (94.6/51.2/29.9/14.4 and 41/54/60/72 — verified against the Summary sheet); action effects are Aid–Bonus Welch differences on all responses; Moral-share effects on the classified sample.

The logic. The round-1 ordering yields an ordinal prediction: story effects should fade as story–game structural similarity falls. They do, monotonically: −5.6, −4.7, −2.4, +0.5 (n.s.) pp on actions; Moral share

−11.7/−14.5/−9.7/−0.4. A Spearman rank correlation across the 4 games is computed in the stats file but deliberately not printed — the paper calls it “an ordinal, exploratory consistency check,” which is the right altitude.

Honest assessment. Four points, LLM-measured x — it proves nothing alone and the text says so. Two prepared answers: (i) **DG-LT’s presence is not cherry-picking**: its story cells were fielded and pre-registered precisely to supply a fourth point (the text says this; the decision log records the placement debate); (ii) the TG null is triply determined (lowest similarity, no Moral shift, no action shift) and the text explicitly refuses to adjudicate *why* the TG is unmoved from this table alone.

7 Equation (7) — BCS (NG: “beliefs about P2?”)

What it is. BCS = **between-category share** of outcome variance within a game×condition cell: $\sum_c q_c(\bar{t}_c - \bar{t})^2 / \text{Var}(t)$ — computationally the R^2 of the action on the three category dummies (the table note says this). Nothing more exotic.

The answer to his margin note. No — beliefs about P2 are deliberately *not* in eq. (7). The representation has two measured components (stated category $\leftrightarrow$ attention; elicited beliefs $\leftrightarrow$ $\mathbf{x}$). Eq. (7) and Table 13 use the **coarsest cut, categories only**, so the BCS is a *lower bound* on representation-explained dispersion. The belief component enters where it belongs: the benchmark regressions, Figure 12’s predictions, and Table 14’s cells. Also note: in the UG control, beliefs are flat across categories (48.2–48.5%), so there is little between-belief variance to add; in the TG beliefs differ by category (22.0/31.5/35.7%), so category dummies already absorb part of it.

Run and answered (2026-07-20, script 20 section A; now a §5.1 footnote at eq. 7). On the belief-nonmissing subsample, with terciles computed within game on the pooled four conditions, refining cells from categories to category×tercile raises the between-cell share by +0.005 to +0.048 across the eight UG/TG condition cells (≤ 0.06 under any tercile convention, section A2): UG Control 0.114→0.146, UG Market 0.067→0.072, TG Control 0.280→0.315, TG Market 0.260→0.303 (Bonus/Aid similar). So the stated category carries nearly all of what the two measured components explain jointly — in the UG because beliefs are flat across categories, in the TG because category dummies already absorb the category-specific beliefs. One nuance if NG compares to Table 13: the categories-only TG Control share is 0.28 on this subsample vs 0.202 on the full classified sample (131 of 584 TG control senders lack a hypothetical belief) — the footnote therefore compares like with like on the subsample. Mechanical caveat: 9 cells vs 3 can only raise an R^2 , and some cells are thin (min $N = 2$ –6 in control/story conditions), so the increments are if anything upper bounds — which strengthens the “categories carry it” reading.

8 Table 13 — Variance Decomposition (the Market row)

Where it lives. 05_paper_v2_new_outputs.py (E3) → heterogeneity_decomposition_main.tex; appendix version has all 16 game×condition rows. Classified responses only; plain variance ratio; **no inference** (point estimates only).

The Market row’s content. DG-KW: BCS collapses 0.663 → 0.136 while SD *rises* 0.204 → 0.263. The arithmetic worth having on the call: between-variance falls 0.028 → 0.009 while within-variance more than quadruples, 0.014 → 0.060. So the headline fact is not “categories stop mattering”; it is that the frame **moves behavior massively within stated categories** (the equal-split anchor dissolves inside every category), on top of moving the mixture. Appendix pattern: same collapse in the other fixed-pie games

(DG-LT .49→.11, UG .11→.06 with SD doubling), while the **TG between-share rises** (.20→.25) — the game where representations and beliefs keep organizing behavior under the frame.

Why the model is comfortable with this: the market lever is κ_d — it changes the *description*, so it can move attention weights within a coarse verbal category, not only move people across categories. The receiver-side decomposition (App. B) and Table 14’s UG-Market row (16%) say the same thing independently. Fig. 9 (distributions) is the picture to point at.

Honest caveats to concede if pressed.

- *Classified-only sample:* unclassified answers are excluded, and they are 17.8% of DG-KW Market (vs ~2% control). Their authors transfer slightly below classified same-condition participants; excluding them plausibly *understates* within-category dispersion, but the direction of bias on BCS is not obvious. **Resolved without forced labels (2026-07-20, script 20 section D):** treating the unclassified answers as their own observed fourth category — no rescued labels involved — leaves every one of the 16 game×condition between-shares within 0.015 of the classified-only value (DG-KW Market 0.136→0.121, its unclassified $N=71$). The Market collapse is not an artifact of the exclusion.
- *Classification noise is plausibly higher under Market* (answers are shorter, 92% of unclassified mention no counterpart): misclassification mechanically attenuates the between-share. Part of 0.66→0.14 may be measurement. The honest line: the within-category movement is corroborated by exhibits that do not use the classifier at all (the distribution figure, the belief shifts), so the collapse is not *only* noise — but we cannot put a number on the split.

9 Figure 13 — Predicted vs Actual TEs on Mean and SD (NG: “why mean? are beliefs changing?”)

Where it lives. Same E3 block → heterogeneity_predicted_actual.png (6 cells plotted; the 8-row table with DG-LT is App. D). Classified only; no inference.

How it works — pure composition, no regression. For each comparison, the category-conditional action distributions (mean m_c and variance v_c , computed on the two conditions *pooled* — the construction is symmetric in the arms) are held fixed; only the category shares w_c vary by arm. Predicted mean = $\sum w_c m_c$; predicted variance = law of total variance of the mixture. Predicted TE = difference of these across arms, against the actual difference; same object for the SD. This is literally the model’s §2.2 display $\Delta E[t] = \sum \Delta q_c \bar{t}_c$ taken to the data, plus its dispersion analogue.

Answers to his two margin questions.

- “Why mean?” Two reasons. (1) The mean panel is the regression-free counterpart of Figure 12’s P1-action panel: same target, weaker machinery — it shows how much of the mean TEs *composition alone* carries (all signs right; market magnitudes understated, ratios 0.28–0.46). (2) It calibrates the reader for the SD panel: same reweighting, same message — “one message read twice.” It is not a duplicate of Figure 12 because Figure 12’s predictions also move beliefs (and, for P2, the fitted P1 action).
- “Are beliefs changing?” Yes — and that is precisely the point. Under the market frame hypothetical beliefs move a lot (+17.2pp UG acceptance, +10.3pp TG return at the reference action); this exercise *deliberately excludes* them, so the market cells are exactly where it understates. The belief channel

is quantified next door: Figure 12’s footnote — categories alone $R^2 = .80$, adding beliefs .87, with beliefs mattering mainly for the TG. NG’s intuition is the paper’s own account of the attenuation.

Honest assessment. Sound and honest as constructed; the caption could state “pooled across the two conditions” for the conditional distributions. Same classified-only caveat as Table 13.

10 Table 14 — Decomposition of TEs into Representation and Behavior Components

Where it lives. `replication_package/code/11_oaxaca.py` → `oaxaca_catbelief.tex`. Classified sample; UG/TG additionally require non-missing hypothetical beliefs. **Inference added 2026-07-22:** bootstrap SEs are now displayed under Δ mean and every level component ($B = 1,000$ participant-level draws within game $\times$ condition strata, seed 20260721; terciles and the full decomposition recomputed in each draw).

How it works. Cells = stated category $\times$ within-game terciles of the hypothetical belief (terciles computed on the two compared conditions) in UG/TG; categories only in the DG (no belief elicitation). Two components: representation = $\sum_c \Delta q_c \cdot \bar{y}_c(\text{ref})$; behavior = $\sum_c q_c(\text{ref}) \cdot \Delta \bar{y}_c$. Three reference schemes: **symmetrized** (average of the two Oaxaca directions — the preregistered headline), **baseline-reference** (Control or Bonus; the column header now reads “Baseline ref.”), **pooled-reference** (all four conditions, with the two-condition tercile edges applied to the pooled sample — exact by construction since the 2026-07-22 revision). Cells present in only one condition contribute their observed mean to the representation component and nothing to behavior.

The findings. Representation carries 39–65% of the TE in the DG and TG cells and about half of the UG story effect. The UG-Market row is the flagged exception: 16% symmetrized — the frame moves behavior *within* cells — and the TG-Aid share is now suppressed (—) in the table because the total is +0.8pp (bootstrap SE 2.0pp) and the share is undefined or degenerate in 44% of bootstrap draws (it was 144% in the version NG read).

Honest assessment — three things to have ready (prereg variants now run, script 20 sections B/C; results in a new footnote at the Table 14 paragraph).

1. **Prereg match — resolved with numbers.** The *method* (symmetrized Shapley/Oaxaca–Blinder) is word-for-word preregistered in both documents. The *cells* are richer than the prereg text: AsPredicted #286187 says “conditional on each representation category” (no terciles); #261370 (DG) words the decomposition over “social proximity levels (representations).” The wording-literal variants are now computed (after first validating the reimplementations: the category-only DG-KW rows reproduce the published $-0.101/39\%$ and $-0.059/65\%$ exactly). On the same samples as Table 14: category-only cells give UG Market 11% (vs 16%), UG Aid 47% (vs 52%), TG Market 35% (vs 59%); DG-KW by social proximity gives Market 36% (vs 39%) and Aid 50% (vs 65%). Same picture throughout, with one informative exception — **TG Market 35% vs 59%: the belief terciles carry roughly 24pp of the representation component in the trust game**, exactly the game where beliefs differ by category. That is now a disclosed footnote, and it is the single best talking point for NG’s “beliefs about P2?”/“are beliefs changing?” notes: yes — and the decomposition can show precisely where.
2. **Convention-dependence where it matters.** UG-Market: 16% symmetrized, $\approx 1\%$ at baseline reference, 40% at the corrected pooled reference. That spread is not a bug — when behavior moves within cells, any Oaxaca split is convention-dependent — but the honest statement is “the representation

share in UG-Market is small under every convention, and the spread across conventions is itself the evidence of within-cell movement.” Don’t let the 16% be quoted alone.

3. **Thin cells — inventoried, and better than feared.** Per-cell N s are now logged (script 20, section C): **no cell is empty in either condition of any comparison, so the empty-cell rule in the notes never binds** (also stated in the new footnote). Thin cells exist and sit where expected — UG $\text{MBC} \times \text{tercile}$ cells run $N=3-7$ in Control/Bonus/Aid (MBC is $\sim 3\%$ of UG answers); TG has three cells under 10 (Moral $\times$ high and SI $\times$ high in Control, SI $\times$ high under Aid). Everything else is comfortably double-digit.

Revision applied 2026-07-22 (AA’s four T14 notes; agreed with SN — bootstrap yes, pooled kept, terciles corrected). The compile NG read is frozen as `main_v2.tex`; his highlighted PDF shows the pre-revision table.

- **Pooled reference corrected.** The old construction cut pooled terciles on all four conditions while the cells used two-condition cuts (14.5–31.9% of participants changed tercile label) and did not close (log-only residual up to $+0.046 = 21\%$ of the UG-Market TE). Now: two-condition edges applied to the pooled sample, behavior = exact complement; $\text{rep} + \text{beh} = \Delta$ mean to machine precision. Displayed column moves ≤ 0.006 : UG-Market $-0.082 \rightarrow -0.088$, TG-Market $+0.088 \rightarrow +0.089$, UG-Aid $-0.018 \rightarrow -0.017$, TG-Aid $+0.005 \rightarrow +0.009$.
- **Bootstrap SEs displayed** (see “Where it lives” above); representation shares stay point estimates; TG-Aid share dashed. The share CIs are wide even where components are cleanly signed (UG-Market share 95% CI [3%, 30%]; TG-Market [42%, 80%]; DG-KW Market [11%, 73%]; TG-Aid [−508%, +810%]) — quote component SEs, not share CIs.
- **“Control ref.” $\rightarrow$ “Baseline ref.”** (the baseline is Bonus in Aid vs Bonus).
- Symmetrized and baseline columns numerically unchanged; all table/figure numbers verified unchanged via aux diff (NG’s highlights and the sent email’s pointers remain valid); `main.tex` 111 pp., 0 undefined references.
- **Open for the call:** AA proposes dropping the two reference columns entirely (symmetrized-only — preregistered, standard); they were NG’s own July-19 request (“alloc media in control” / “tutto il pooled sample”), so decide with NG. If dropped, keep the UG-Market convention-spread point ($\approx 1\%$ baseline / 16% symmetrized / 40% pooled = within-cell movement) as a prose or note sentence.